

# Supervised Learning for NLP I

HSS 510 / DS 518: NLP for HSS

Taegyeon Kim

Mar 18, 2025

# Agenda

## Things to be covered

- Overview of supervised learning
- Step 1: Building a labeled data set
- Step 2: Training models
- Step 3: Evaluating performance
- Guided labeling: text classification with movie reviews data in Python

# Supervised Learning

## Supervised vs. Unsupervised

	Supervised	Unsupervised
Objective	Trained on a labeled data to learn a mapping from input to output	Find patterns or structures within data without labels
Outcome	Pre-defined categories	Not quite pre-defined
Model evaluation	Explicit metrics such as accuracy, precision, recall, or MSE	Can involve qualitative assessment
Examples	Classification/regression for texts	Topic models

# Supervised Learning

## Regression vs. Classification

- Regression
  - The outcome of interest is continuous or ordered (beyond binary)
  - E.g., OLS regression (+ non-linear regression algorithms such as random forest regression)
- Classification
  - The outcome is a value in an unordered set (i.e. categories)
- The two approaches share the broad principles of supervised learning and can be adapted

# Supervised Learning

We will focus on text classification with supervised learning

- Goal
  - To classify documents into pre-defined categories
  - E.g., sentiment of comments (e.g., positive-negative), stance on issues (e.g., for-against-neutral), etc.
- We need
  - Labeled data set (for training and testing)
  - Model (algorithm) that maps texts to labels
  - Evaluation approaches: performance metrics, cross-validation, etc.

# Supervised Learning

## Evolution of text classification

- Dictionary methods
  - Based on counting/weighting of relevant keywords
  - Readily available and fast  $\leftrightarrow$  sub-optimal performance
  - If you want to quick, preliminary analysis for concepts supported by an existing dictionary, it can be helpful
  - E.g., [LIWC](#), [VADER](#), [Moral Foundations Dictionaries](#), etc.

# Supervised Learning

## Evolution of text classification (cont'd)

- Traditional ML algorithms
  - This approach laid the groundwork for many foundational concepts in text classification
  - Classifiers (models) are trained to learn the relationships between texts and labels (i.e., classes)
  - So this requires training (labeled) data (pairs of a text and a label)
  - (On average) more training data, higher performance
  - E.g., Logistic regression, random forest, SVM, deep neural networks

# Supervised Learning

## Evolution of text classification (cont'd)

- Fine-tuning representation models (e.g., BERT family)
  - These models are pre-trained with massive amounts of text data
  - Given a text, representation models encode it into a vector (an array of numbers) that captures its meaning
  - We can fine-tune such a model for classification with potentially higher performance
  - They tend to achieve higher performance than the traditional ML approaches
  - [Week 13](#) will cover this approach



# Supervised Learning

## Evolution of text classification (cont'd)

- Prompting generative models (e.g., GPT family)
  - Like representation models, these models are pre-trained on vast amounts of text data, often even more extensively
  - Generative models are designed to generate text outputs given input text prompts
  - We can prompt such a model with unlabeled texts to generate labels
  - Still requires labeled data (not for training but for evaluating performance)
  - It is also possible to “fine-tune” these models
  - [Weeks 14–15](#) will cover this

# Overview of Process

## Overall process

- Step 1: building a labeled data set
- Step 2: training model(s)
- Step 3: evaluating performance

# Overview of Process

## Overall process

- **Step 1: building a labeled data set**
  - **Step 2: training model(s)**
  - **Step 3: evaluating performance**
- \* Steps 1 and 3 apply to all of the classification approaches**

# Overview of Process

## Step 1: build a labeled data set

- Label texts following systematic labeling guidelines and check inter-coder reliability
- This ( $C$ ) will serve as “ground-truth” or “gold standards”
- $C$  is used to training (building a model) and validation (evaluating performance)

# Overview of Process

Step 1: build a labeled data set

Doc number	Text	y
1	This is great!	0
2	%@% off!	1
...		
9999	This is sick	0
10000	Love BTS <3	0

# Overview of Process

Step 1: build a labeled data set

Doc number	Text	<i>y</i>
1	This is great!	0
2	%@% off!	1
...		
9999	This is sick	0
10000	Love BTS <3	0

# Overview of Process

## Step 2: train models

- Randomly split  $C$  into a training set ( $C_{train}$ ) and a test set ( $C_{test}$ )
  - Typically,  $C_{train} : C_{test} = 7:3$  or  $8:2$
  - E.g., identifying YouTube comments containing hate speech
    - $C$ : 10,000 comments labeled for the presence of hate speech
    - $C_{training}$ : 8,000 comments for training
    - $C_{test}$ : 2,000 comments for test
- Generate  $X_{train}$  (feature matrix) from  $C_{train}$ 
  - E.g., count vectors, TF-IDF, or embeddings

# Overview of Process

Step 2: train models

Index	y	Token 1	Token 2	...	Token V-1	Token V
1	0	3	1.4	...	1.7	6
2	1	-0.8	6.4	...	5.7	-1.6
...						
7999	0	-2.8	0.9	...	3.3	-0.6
8000	0	3.7	1.4	...	5.7	-5.8



# Overview of Process

## Step 2: train models

- Choose a model  $F$  (e.g., logistic regression) and learn model parameters  $\beta$  (e.g., an array of coefficients)
  - The model provides a mapping between  $X_{train}$  and  $y_{train}$
- Loss (cost) function: measures how much model predictions ( $\hat{y}_{train}$ ) differ from the true labels ( $y_{train}$ )
  - $\hat{y}_{train} = F(\hat{\beta} * X_{train})$
  - $\beta$  is estimated in a way that minimizes the difference between the two
- As a result, we get a classifier:  $\hat{y} = F(\hat{\beta} * X)$

# Overview of Process

Step 2: train models

Index	$y$	$\hat{y}$	Token 1	Token 2	...	Token V-1	Token V
1	0	1	3	1.4	...	1.7	6
2	1	1	-0.8	6.4	...	5.7	-1.6
...							
7999	0	0	-2.8	0.9	...	3.3	-0.6
8000	0	0	3.7	1.4	...	5.7	-5.8

# Overview of Process

## Step 3: evaluate performance

- We held out another labeled set  $C_{test}$  ( $n = 2,000$ ) (**why?**)
- Use the classifier  $F(\hat{\beta} * X)$  to generate predictions  $\hat{y}_{test}$
- Compare the predictions  $\hat{y}_{test}$  and the true labels  $y_{test}$
- Performance metrics include accuracy, precision, recall, etc.
- (Then use the classifier for unlabeled data)

# Overview of Process

Step 3: evaluate performance

Index	$y$	$\hat{y}$	Token 1	Token 2	...	Token V-1	Token V
1	1	1	3.12	1.99	...	5.77	0.36
2	1	0	-0.8	1.14	...	9.71	-1.66
...							
1999	0	0	-2.11	0.95	...	1.23	-0.62
2000	0	0	3.71	1.48	...	1.7	-5.84

# Bias, Variance, and Overfitting




## Bias

- The degree to which the model's predictions deviate from the true labels in a systematic manner
- A model with high bias make predictions that are consistently off-target


## Variance

- The degree to which the model generalizes to different data
- High variance means low generalizability


# Bias, Variance, and Overfitting




# Bias, Variance, and Overfitting



**overfitting**



**underfitting**




**Good balance**


# Bias, Variance, and Overfitting

## Overfitting and underfitting


- If a model learns the training data “too well” (low bias), it can lead to overfit
- This happens when the model mistakes noise for signal
- The model would not generalize to the test set (high variance)



**overfitting**



**underfitting**



**Good balance**



# Bias, Variance, and Overfitting

## Training-test split

- A minimal measure to prevent overfitting
- The primary goal here is to make our model as generalizable as possible
- “Generalizable” means being able to perform well on unseen documents (other than the documents the model was trained on)
- When a model learns the noise or random fluctuations in the training set, this typically results in a model that performs poorly on new the test set

# Step 1: Building a Labeled Data Set

In HSS and many other fields, text classification is often used to “measure” concepts using texts

- Some concepts are already available in the form of labeled texts: e.g., news categories
- In many cases, however, concepts we are trying to measure using texts need to be measured from scratch
- The process involves definition/conceptualization and manual labeling

# Step 1: Building a Labeled Data Set

How do we build a labeled data set?

- A form of manual content analysis
- Here human labels are often considered “ground truth” or “gold standard”
- Manually labeled data are used for training (train set) and evaluation (test set)

# Step 1: Building a Labeled Data Set

How do we build a labeled data set?

- Expert labeling
  - In many projects, a few domain experts work on labeling
    - E.g., a researcher + two RAs
  - Define the concept and draft coding guidelines
  - Annotators should be trained to learn the concept and related guidelines

# Step 1: Building a Labeled Data Set

How do we build a labeled data set? (cont'd)

- Crowd-sourced labeling
  - “Wisdom of crowd”: aggregated judgments of (online) non-experts converge to judgments of experts at much lower cost ([Benoit et al, 2016](#))
  - Difficult to educate annotators on sophisticated tasks
  - Inductive measurement based on a relatively loose conceptualization

# Step 1: Building a Labeled Data Set

## Expert labeling vs. Crowd Sourcing

- Deductive vs. inductive
- Degree of training
- Scalability (cost)

# Step 1: Building a Labeled Data Set

## Selecting texts for manual labeling

- The selected texts should be representative of the entire corpus.
- These texts are used to train and test your model
- Mismatches between the chosen texts and the texts the model is applied to can reduce generalizability
- E.g., changes in anti-vaccine discourse throughout 2020
  - After mid-2020, anti-vaccine discourse increasingly focused on government mandates
  - If the model is trained only on texts from the first two months of 2020, it may fail to capture these later shifts

# Step 1: Building a Labeled Data Set

## Iterative process

- The process of building a labeled set does not often take place at once but in an iterative process
- In many cases, specifying comprehensive annotation guidelines *ex ante* can be challenging
- Preliminary labeling rules are written and applied to an initial set of texts → annotators identify disagreements in their labels and ambiguities in the rules → revision of the rule → manual annotation of another set of texts ...



# Step 1: Building a Labeled Data Set

## Dealing with subjectivity

- Many concepts in humanities and social sciences tend to be subjective
- This is, from the beginning, why 1) careful conceptualization and 2) writing clear labeling rules, and 3) training coders are extremely important
- Once your final labels are built, it is helpful to assess reliability between annotators (inter-coder reliability)
  - Krippendorff's  $\alpha$ , Cohen's  $\kappa$  (alternatives include Pearson's  $r$ , Spearman's  $\rho$ )
  - Relevant tools: [krippendorff](#) in Python or [irr](#) in R

# Step 1: Building a Labeled Data Set

Who are the annotators?

- Expert coding
  - Academics/students ([Javdani and Chang 2023](#))
- Crowdsourcing
  - Skewed distribution of worked hours ([Difallah et al. 2018](#))
  - Inattentive workers ([Peyton et al. 2022](#); [Ternovski 2022](#))
  - LLM-based responses ([Veselovsky et al. 2023](#))
  - Demographic characteristics ([Al Kuwatly et al. 2020](#))

## Step 2: Training Models

We the need to generate features

- Features mean numerical representations of text data that models can process
- $C_{train} \rightarrow X_{train}$
- Options include count vectors, TF-IDF vectors, word/document embeddings, etc.
- Note, as we will see in later weeks, we hardly need manual feature extraction when we fine-tune or prompt pre-trained models

## Step 2: Training Models

We the need to generate features (cont'd)

Index	Token 1	Token 2	...	Token V-1	Token V
1	3	1.4	...	1.7	6
2	-0.8	6.4	...	5.7	-1.6
...					
7999	-2.8	0.9	...	3.3	-0.6
8000	3.7	1.4	...	5.7	-5.8

## Step 2: Training Models

So far we have:

- Built a labeled data set
- Generated a feature matrix
- This means that we have the outcome ( $y$ ) and features ( $X_{train}$ )

Now we will:

- Select a model  $F$
- Learn model parameters  $\beta$  to build a classifier ( $\hat{y} = F(\hat{\beta} * X)$ )

# Step 2: Training Models

Numerous algorithms

- Logistic regression
- Naive Bayes
- Support vector machine
- Tree-based models (decision tree, random forest, XGBoost, etc.)
- Neural networks
- Etc.

# Step 2: Training Models

## Logistic regression

- Used to classify a document into binary categories
  - For more than two categories, multinomial logistic regression
- One of the most useful analytics tools in science (not just NLP/ML)
- The baseline supervised learning algorithm for classification
- Forms the basis of neural networks

# Step 2: Training Models

## Components of classification with logistic regression

- Features
  - A document is represented as a vector of features  $\vec{x} = [x_1, \dots, x_n]$
- A classification function ( $F$ )
  - $p(y = 1|x)$  is computed for each document given the feature vector ( $\vec{x}$ ) and the parameters ( $\beta: \vec{w}$  and  $b$ )
  - The sigmoid function transforms  $p(y = 1|x)$  into a value between 0 & 1
- Loss function (measures how model prediction  $\hat{y}$  is different from  $y$  during training) and an algorithm for optimizing it (gradient descent)



## Step 2: Training Models

How does logistic regression compute predicted probabilities?

- $p(y = 1|x)$ 
  - The probability of  $y = 1$  given a feature vector  $x \vec{=} [x_1, \dots, x_n]$
  - E.g., for a simple count vector, it would be # of times each token appears in the document
- Logistic regression learns  $\beta$ , a vector of coefficients
  - A bias term  $b$ : a single number (a.k.a. intercept)
  - Weights  $w \vec{=} [w_1, \dots, w_n]$ 
    - E.g., tokens signaling hateful intention would get high weights
  - With  $b$ ,  $w \vec{}$ , and  $x \vec{}$ , we compute  $z = (\sum_{i=1}^n w_i x_i) + b$


## Step 2: Training Models

How does logistic regression compute predicted probabilities?

- $z = (\sum_{i=1}^n w_i x_i) + b = \vec{w} \cdot \vec{x} + b$

## Step 2: Training Models

Sigmoid function



## Step 2: Training Models

How does logistic regression compute predicted probabilities?

- $z = (\sum_{i=1}^n w_i x_i) + b = \vec{w} \cdot \vec{x} + b$

- $\sigma(z) = \frac{1}{1+e^{-z}} = \frac{1}{1+\exp(-z)}$

## Step 2: Training Models

How does logistic regression compute predicted probabilities?

- $p(y = 1|x) = \sigma(w^{\vec{}} \cdot x^{\vec{}} + b)$
- $p(y = 0|x) = 1 - \sigma(w^{\vec{}} \cdot x^{\vec{}} + b)$

## Step 2: Training Models

How do predicted probabilities turn into binary labels ( $\hat{y}$ )?

$$\begin{cases} 1 & \text{if } P(y = 1|x) > 0.5 \\ 0 & \text{otherwise} \end{cases}$$

## Step 2: Training Models

E.g., sentiment classification from movie reviews

*“It’s hokey. There are virtually no surprises, and the writing is second-rate. So why was it so enjoyable? For one thing, the cast is great, Another nice touch is the music. I was overcome with the urge to get off the couch and start dancing? It sucked me in, and it’ll do the same to you.”*

## Step 2: Training Models

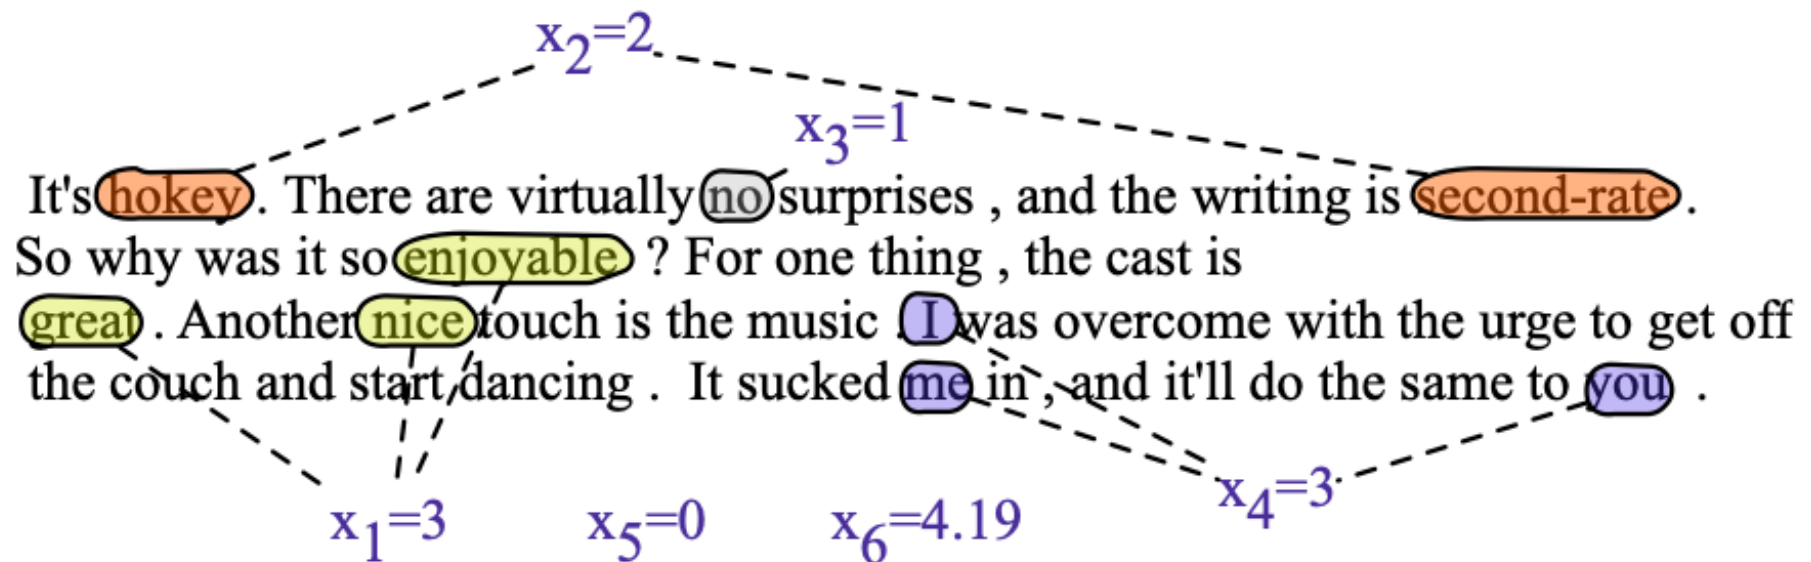
Sentiment classification from movie reviews

Var	Definition
$x_1$	count(positive lexicon) $\in$ doc)
$x_2$	count(negative lexicon) $\in$ doc)
$x_3$	$\begin{cases} 1 & \text{if "no"} \in \text{doc} \\ 0 & \text{otherwise} \end{cases}$
$x_4$	count(1st and 2nd pronouns $\in$ doc)
$x_5$	$\begin{cases} 1 & \text{if "!"} \in \text{doc} \\ 0 & \text{otherwise} \end{cases}$
$x_6$	log(word count of doc)



## Step 2: Training Models

Sentiment classification from movie reviews



## Step 2: Training Models

How are the parameters, weights  $\vec{w}$  and bias  $b$ , learned?

- Goal: learn  $\vec{w}$  and  $b$  that make  $\hat{y}_{train}^i$  for each training observation as “close” as possible to  $y_{train}^i$  (the true label)
- Two components for estimation
  - Metric for “closeness”: loss/cost function (e.g., cross entropy loss)
  - Optimization algorithm: (stochastic) gradient descent

## Step 2: Training Models

Learn parameters that maximize the chance of getting the correct label (Conditional Maximum Likelihood Estimation)

- $p(y|x) = \hat{y}^y * (1 - \hat{y})^{1-y}$
- If  $y = 1$ ,  $p(y|x) = \hat{y}$ 
  - The higher  $\hat{y}$  is, the better the classifier
- If  $y = 0$ ,  $p(y|x) = (1 - \hat{y})$ 
  - The higher  $(1 - \hat{y})$  is (the lower  $\hat{y}$  is), the better the classifier

## Step 2: Training Models

From  $p(y|x)$ , we derive loss (cost) function

- $p(y|x) = \hat{y}^y * (1 - \hat{y})^{1-y}$ 
  - $\text{Log}(p(y|x)) = \text{Log}(\hat{y}^y * (1 - \hat{y})^{1-y})$
  - $\text{Log}(p(y|x)) = y\log\hat{y} + (1-y)\log(1-\hat{y})$

We can turn it into the cross entropy loss function (that should be minimized)

- $L_{CE} = -[y\log\hat{y} + (1-y)\log(1-\hat{y})]$

## Step 2: Training Models

### Gradient descent

- With the cross entropy loss function, we have a metric for discrepancies
- Gradient descent is an algorithm used to find the optimal set of weights (and bias) that minimizes the discrepancies, averaged over all observations

$$\rightarrow \hat{\theta} = \underset{\theta}{\operatorname{argmin}} \frac{1}{m} \sum_{i=1}^m L_{\text{CE}}(f(x^{(i)}; \theta), y^{(i)})$$

## Step 2: Training Models


Gradient descent finds a minimum of a function

- Figures out in which direction the function's slope rises most steeply
- Then move in the opposite direction
- Stops when it reaches the minimum

# Step 2: Training Models

## Gradient descent


- 1-dimensional illustration



# Step 2: Training Models

## Gradient descent

- Updating  $w$  with learning rate ( $\eta$ ):  $w^{t+1} = w^t - \eta \frac{d}{dw} L(f(x; w), y)$





## Step 4: Evaluate Performance

We have

- Manually labeled documents
- Split them into  $C_{train}$  (training set) and  $C_{test}$  (test set)
- Trained a classifier on  $C_{train}$  (with  $y_{train}$  and  $X_{train}$ )  $\rightarrow F(\hat{\beta}^* X)$

Now we need to evaluate its performance on  $C_{test}$

- We compare  $\hat{y}_{test}$  (predicted labels) against  $y_{test}$  (true labels)

# Step 4: Evaluate Performance

## Performance metrics

- Accuracy: the proportion of all predictions (both positive and negative) that the model got right
- Precision: the proportion of positive predictions that were actually correct
- Recall: the proportion of actual positives that were correctly predicted
- F-1: the harmonic (as opposed to arithmetic) mean of precision and recall

# Step 4: Evaluate Performance

Confusion matrix: predictions against true labels

		True condition	
		Positive	Negative
Prediction	Positive	True Positive	False Positive (Type I error)
	Negative	False Negative (Type II error)	True Negative

## Step 4: Evaluate Performance

Accuracy:  $\frac{TP+TN}{TP+TN+FP+FN}$

		True condition	
		Positive	Negative
Prediction	Positive	True Positive	False Positive (Type I error)
	Negative	False Negative (Type II error)	True Negative

## Step 4: Evaluate Performance

Precision:  $\frac{TP}{TP+FP}$

		True condition	
		Positive	Negative
Prediction	Positive	True Positive	False Positive (Type I error)
	Negative	False Negative (Type II error)	True Negative

## Step 4: Evaluate Performance

Recall:  $\frac{TP}{TP+FN}$

		True condition	
		Positive	Negative
Prediction	Positive	True Positive	False Positive (Type I error)
	Negative	False Negative (Type II error)	True Negative

## Step 4: Evaluate Performance

$$F-1: (2 \times \textit{precision} \times \textit{recall}) / (\textit{precision} + \textit{recall})$$

# Step 4: Evaluate Performance

## Reminders

- Random train-test split:  $C_{train}$ ,  $C_{test}$ 
  - E.g., 10,000 comments labeled for hate speech into 8,000 and 1,000
- Our classifier learned **parameters**, maximizing performance on  $C_{train}$  and evaluating it on  $C_{test}$




# Step 4: Evaluate Performance

## Parameters vs. hyper-parameters

- Parameter
  - Learned (estimated) from data (internal to the model)
  - E.g., logistic regression weights/coefficients ( $= \beta$ )
- Hyper-parameters
  - Defines the model structure itself (not internal to the model)
  - E.g., the size of a regularization term in logistic regression, the number of layers or learning rate in neural networks, etc.


# Step 4: Evaluate Performance

Shallow neural network



Hand-designed feature extraction

Deep neural network



Learn a *feature hierarchy* all the way from input to output data

## Step 4: Evaluate Performance

Hyper-parameters influence model performance, and we want to “tune” them

- With different hyper-parameter values, we could fit a model configured with each value on  $C_{train}$  and evaluate performance on  $C_{test}$
- E.g., regularization strength  $\lambda$  in logistic regression
  - Train different models with different  $\lambda$  values on  $C_{train}$
  - Evaluate on  $C_{test}$
  - Pick the best performing model

## Step 4: Evaluate Performance

What could go wrong?

- In comparing different models (different hyperparameter values), we might overfit on  $C_{train}$
- By repeatedly using the specific train-test split, our comparison can be affected by the specific characteristics of the split

# Step 4: Evaluate Performance

## Validation set

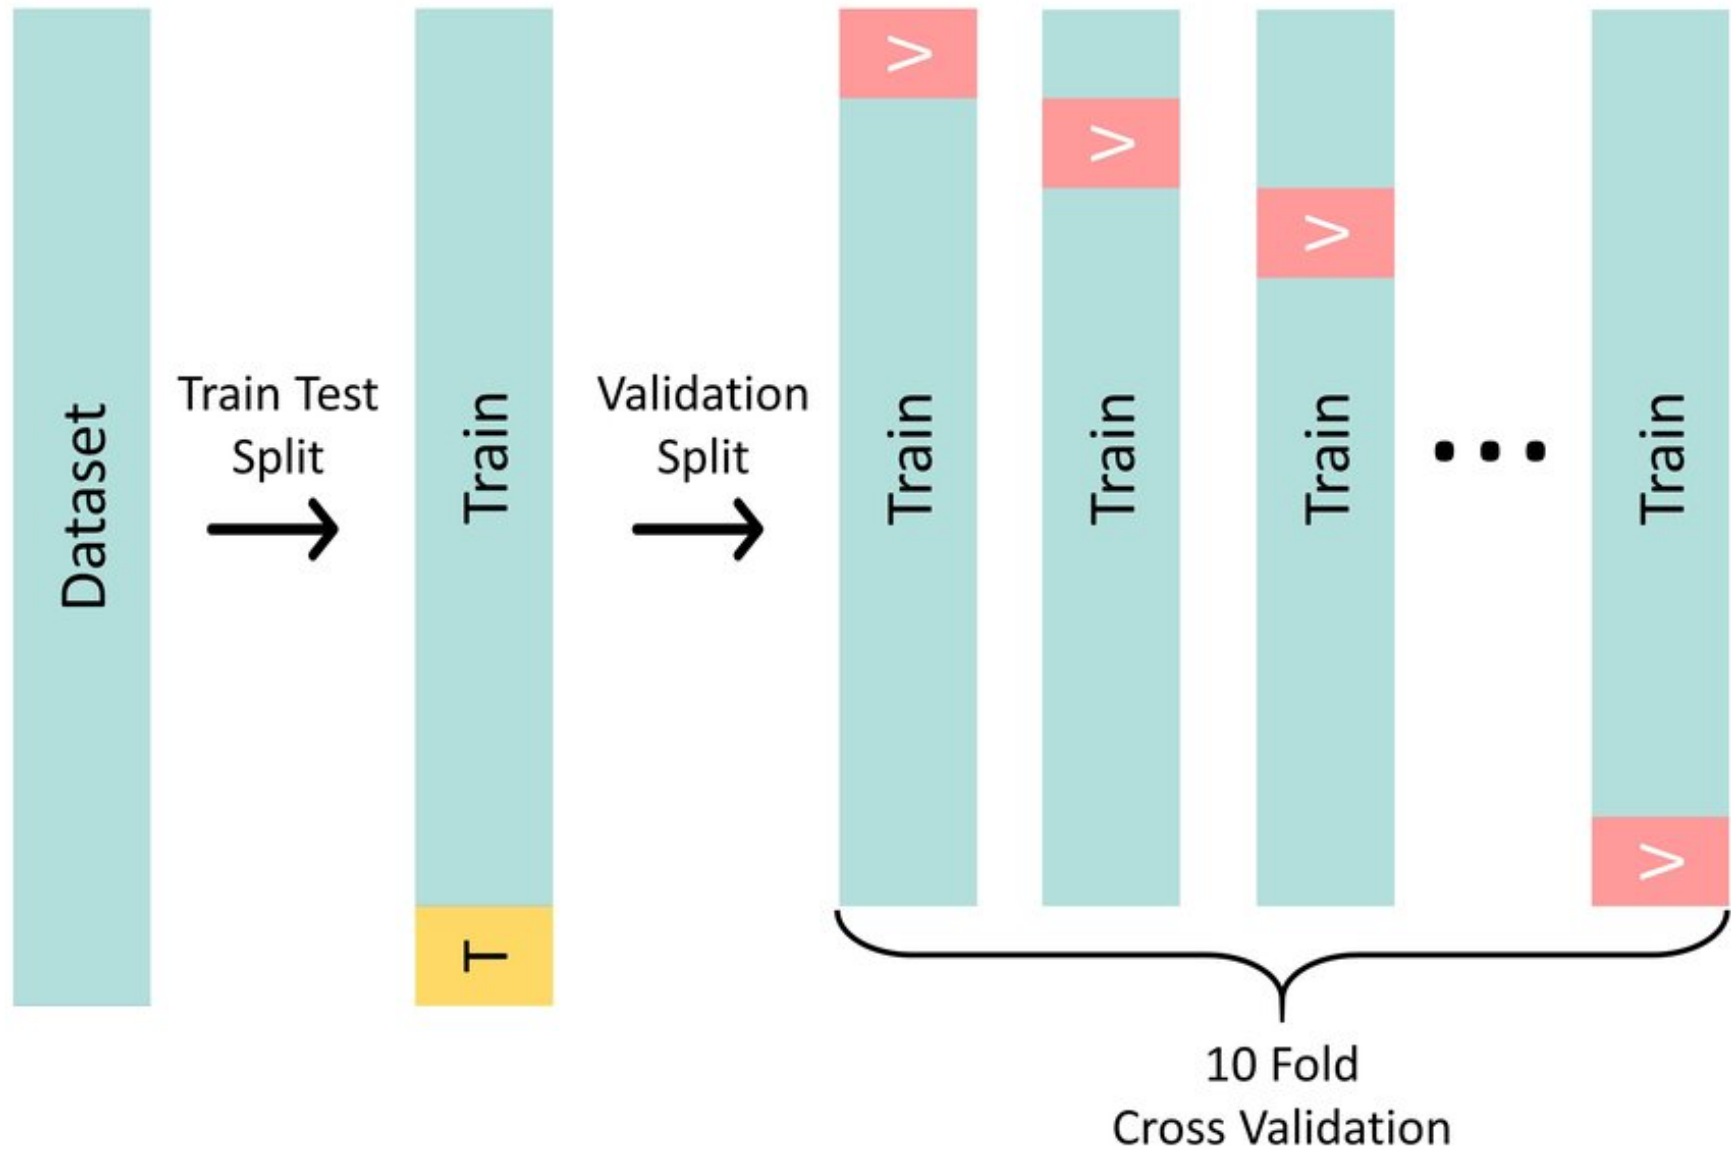
- We split the labeled data set into a training set, a “validation set”, and a test set:  $C_{train}$ ,  $C_{validation}$ , and  $C_{test}$
- Train a model on  $C_{train}$  and see how it performs on  $C_{validation}$
- Repeat this step for multiple hyper-parameter configurations
- Pick the best-performing configuration
- Train a final model based on the configuration using  $C_{train} + C_{validation}$
- Evaluate on  $C_{test}$

## Step 4: Evaluate Performance

*K*-fold cross-validation

- Randomly split  $C_{train}$  into  $K$  equal parts or “folds” (commonly 5 or 10)
- For each iteration
  - Treat one fold as the “validation set”
  - Train your model on the remaining  $K - 1$  folds
  - Evaluate performance on the validation set kept aside
- After cycling through all iterations
  - Aggregate the performance metrics obtained from each iteration
  - Choose the classifier with the highest cross-validated performance
  - This step may involve not just hyper-parameter tuning but also things like feature representation, etc.
- (Re)train the chosen best classifier on  $C_{train}$  (all  $K$  folds combined) and evaluate on  $C_{test}$

## Step 4: Evaluate Performance





# Summary

Supervised text classification provides a highly useful tool to assign labels to texts

- Be aware of the principles of building a labeled data set
  - Conceptualization, inter-coder reliability, annotator bias, etc.
- Validate, validate, and validate!
- Choose appropriate evaluation metrics

# Guided Coding

Text classification with movie reviews data ([link](#))